

# Data Manipulation Statements

- SELECT
- INSERT
- UPDATE
- DELETE
- Refer to documentation for others

The above keywords are used for the so-call "CRUD" operations on data:

**C**reate, **R**ead, **U**ppdate, **D**elete

# Data Definition Statements

- CREATE
  - TABLE
  - VIEW
  - ...
- ALTER
  - TABLE
  - VIEW
  - ...
- DROP
  - TABLE
  - VIEW
  - ...

# Data Definition Statements

- Data definition statements manipulate database objects such as:
  - Tables
  - Views
  - Indexes
  - Triggers
  - Functions
  - Procedures
- MySQL documentation:  
<https://dev.mysql.com/doc/refman/8.0/en/sql-data-definition-statements.html>

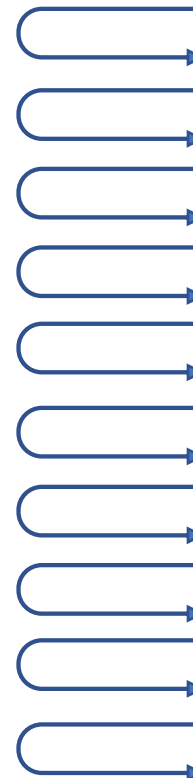
# Indexes

- A structure associated with a table that can speed up data access.
- An index is defined on one or more columns on a table.
- Without an index, this statement would perform a **full table scan**:

```
SELECT * FROM city WHERE CountryCode = 'IND';
```

- **Full table scan:** A scan of table data in which the database sequentially reads all rows from a table and filters out those that do not meet the selection criteria.

world database  
**city** full table scan



CountryCode	Name
AFG	Qandahar
AFG	Herat
...	...
COL	Medellin
...	...
IND	<b>Delhi</b>
IND	<b>Chennai</b>
...	...
NGA	Ibadan
NGA	Kano
...	...
...	...

# world database **city** table

```
CREATE TABLE `city` (  
  `ID` int NOT NULL AUTO_INCREMENT,  
  `Name` char(35) NOT NULL DEFAULT '',  
  ...  
  PRIMARY KEY (`ID`),  
  KEY `CountryCode` (`CountryCode`),  
  CONSTRAINT `city_ibfk_1` FOREIGN KEY (`CountryCode`)  
REFERENCES `country` (`Code`)  
);
```

city search by index

city index

CountryCode
AFG
...
COL
...
IND
...
NGA

city table

CountryCode	Name
AFG	Qandahar
AFG	Herat
...	...
COL	Medellin
...	...
IND	Delhi
IND	Chennai
...	...
NGA	Ibadan
NGA	Kano
...	...
...	...

# Multiple column indexes

```
CREATE TABLE test (  
    id INT NOT NULL,  
    last_name CHAR(30) NOT NULL,  
    first_name CHAR(30) NOT NULL,  
    PRIMARY KEY (id),  
    INDEX name (last_name,first_name)  
);
```

The index can be used for lookups in queries that specify values in a known range for combinations of last\_name and first\_name values.

It can also be used for queries that specify just a last\_name value because that column is a **leftmost prefix** of the index.




composite index

last_name, first_name
Cooper, Andy
Jones, John
Kaushik, Sameer
...

table

id	first_name	last_name
4	Andy	Cooper
5	John	Jones
...	...	...
12	Sameer	Kaushik
...	...	...



Queries using index:

```
SELECT * FROM test WHERE last_name='Jones';
```

```
SELECT * FROM test WHERE last_name='Jones' AND first_name='John';
```

Queries that will not use index:

```
SELECT * FROM test WHERE first_name='John';
```

```
SELECT * FROM test WHERE last_name='Jones' OR first_name='John';
```

# Indexes

Indexes are used to find rows with specific column values quickly. They are defined on one or more columns.

MySQL Index keywords:

- PRIMARY KEY
- UNIQUE
- INDEX
- FULLTEXT: for text-based columns

Primary and foreign keys have indexes, by default.

# Indexes

## Benefits:

- Find rows with specific values quickly
- Avoid full table scans
- Eliminates rows from consideration
- Create indexes on foreign key columns to makes joins more efficient
- Sorting

# B-tree

A tree data structure that is popular for use in database indexes.

- The structure is kept sorted at all times, enabling fast lookup for exact matches (equals operator) and ranges (for example, greater than, less than, and BETWEEN operators).

# B-tree index

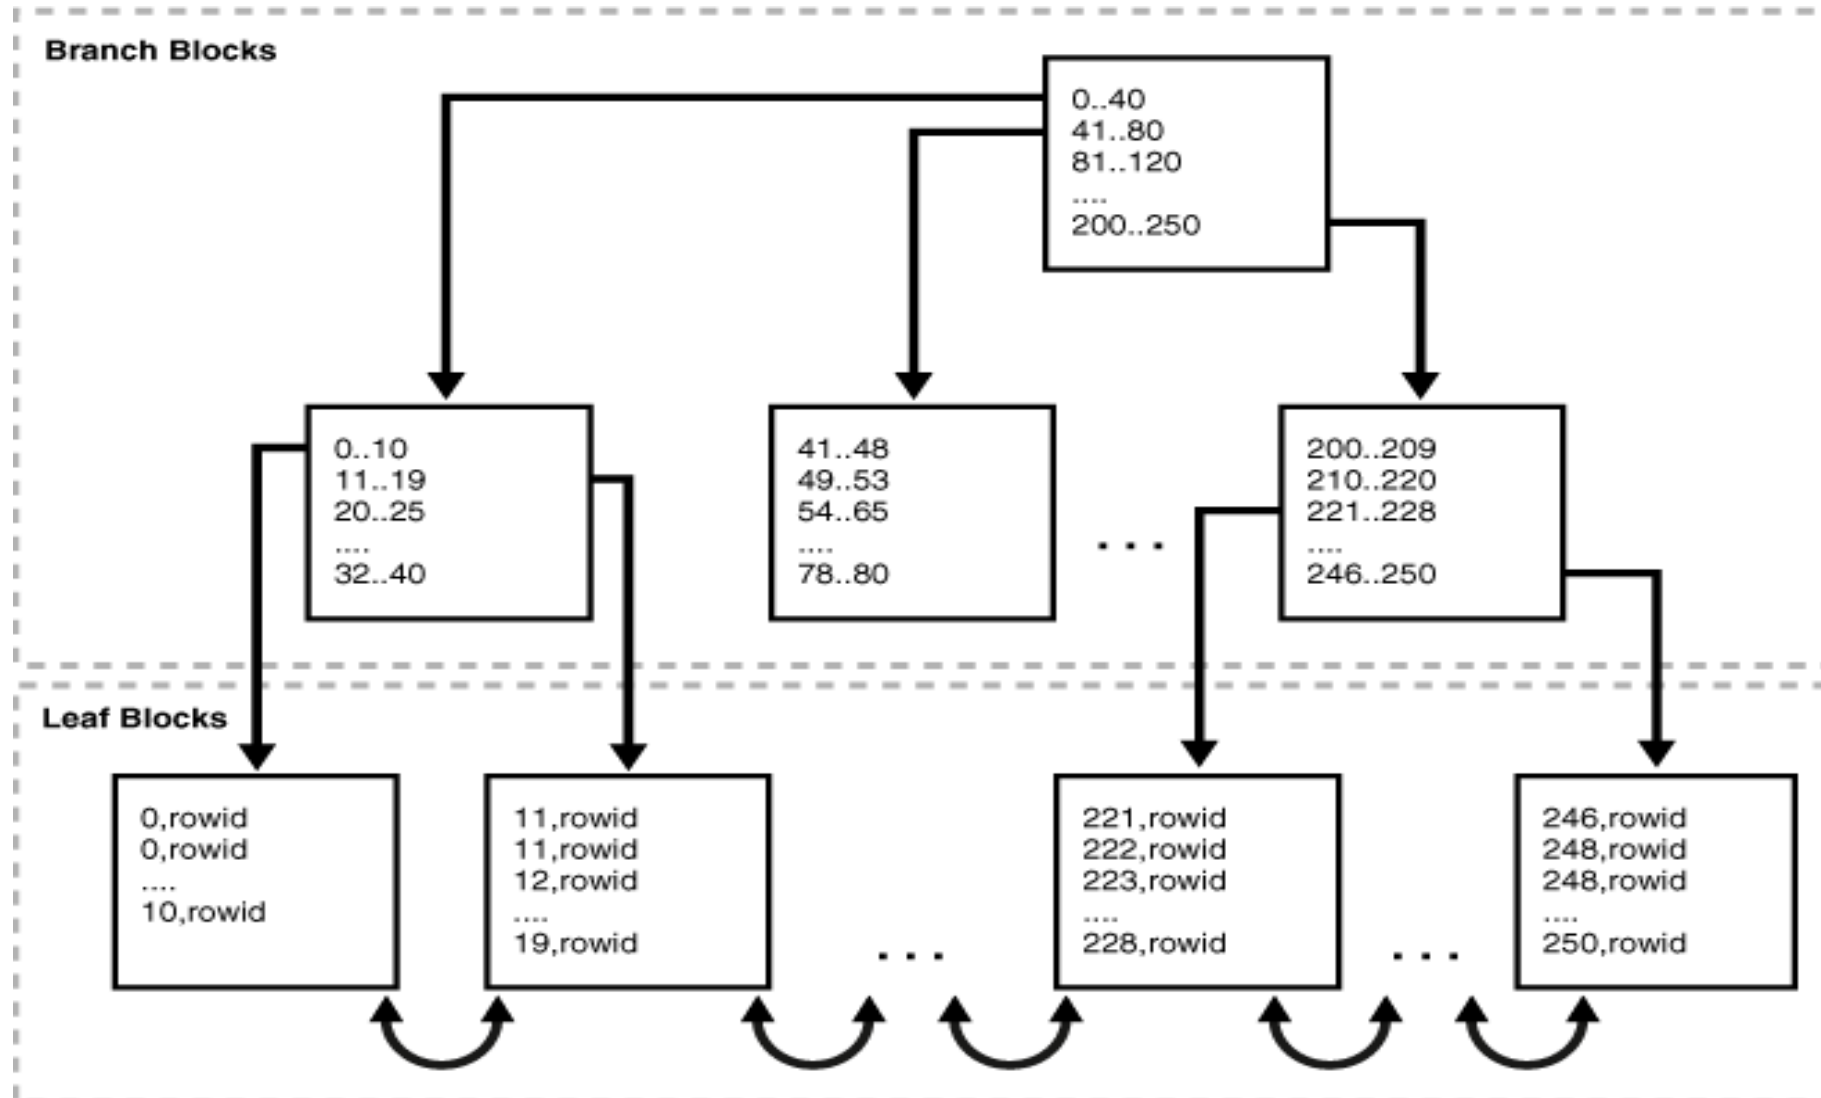


Image source:  
[https://docs.oracle.com/cd/E11882\\_01/server.112/e40540/indexiot.htm#CNCPT721](https://docs.oracle.com/cd/E11882_01/server.112/e40540/indexiot.htm#CNCPT721)

# Definitions

- rowid: A unique address for a row in a database
- Branch blocks: Used for searching for values
- Leaf block: Stores the indexed value and *rowid* to locate the actual row in the referenced table.
- A B-tree index is balanced: all leaf indexes have the same depth. Retrieval of any record in the index takes approximately the same time.

# Index tradeoffs

- Indexes take up extra space.
- When an index is defined on a table, inserts, updates and deletes take longer as both the table and the index are modified.
- Use indexes only when necessary to speed up queries!

# MySQL avoids using indexes in these cases

- The table is so small it is faster to do a full table scan than a key (index) lookup.
- There are no usable restrictions in the ON or WHERE clause for indexed columns.
- You are comparing indexed columns with constant values and MySQL has calculated (based on the index tree) that the constants cover too large a part of the table.
- You are using a key with low cardinality (many rows match the key value) through another column.



View all indexes on 'sakila'

```
SELECT DISTINCT  
    TABLE_NAME,  
    INDEX_NAME  
FROM INFORMATION_SCHEMA.STATISTICS  
WHERE TABLE_SCHEMA = 'sakila';
```

# Readings on indexes

- <https://dev.mysql.com/doc/refman/8.0/en/mysql-indexes.html>
- <https://dev.mysql.com/doc/refman/8.0/en/index-btree-hash.html#btree-index-characteristics>